



Improving performance through technology; delivering measurable results.

# Performance goes up. The numbers prove it.

We find the bottleneck, measure it, improve it with technology, then measure it again the same way. You see the result in your own numbers.

$\Delta$  = performance after - performance before



# Most AI work starts with a tool. We start with work that is stuck.

TOOL - FIRST

BOTTLENECK - FIRST

|                                              |                                                     |
|----------------------------------------------|-----------------------------------------------------|
| Starts with a technology and looks for a use | Starts with a queue of work that waits              |
| "It feels faster"                            | Minutes per unit, before and after                  |
| The vendor reports the result                | The process owner signs the baseline and the result |
| Promises "10x" before looking                | Sets the target after the baseline is measured      |
| Leaves after launch                          | Keeps measuring, then moves to the next bottleneck  |
| Sells licences and hours                     | Sells a measured change, and can put its fee on it  |

The discipline is old: Lean, Six Sigma's Define–Measure–Analyze–Improve–Control, and the Theory of Constraints. We apply it to AI work, where almost nobody measures.

## Six places where the difference is easy to see.

Each tab shows one costly problem, the startup whose success proves it is real, and a before-and-after measurement from a fictional company with realistic numbers.



HUMAN RESOURCES · FICTIONAL EXAMPLE

## Finding star candidates in hours, not days

Group recruiting team. 1,200 resumes a month arrive by email and job boards.

PRIMARY METRIC · OPERATOR MINUTES PER RESUME



|                                      |                    |
|--------------------------------------|--------------------|
| Time to first review                 | 3.5 days → 4 hours |
| Staff-hours per month                | 120 h → 6 h        |
| Top candidates in the tool's top 20% | not measured → 78% |
| GUARDRAIL Field accuracy in the ATS  | 96% → 98% held     |

### WHAT CHANGED

The tool reads each resume, fills the ATS record, and orders new candidates so recruiters open the likeliest matches first. People make every decision.

### WHY THIS PROBLEM IS REAL

The HR problem atlas lists 31 startups built on resume screening and matching alone.

Fictional company. The figures are illustrative and anchored to the real benchmarks named on this tab.



PRODUCTION & PLANNING · FICTIONAL EXAMPLE

## Predicting machine failures before they stop the line

Beverage bottling plant. Two filling lines, 12,000 bottles an hour each, two shifts.

PRIMARY METRIC · UNPLANNED DOWNTIME PER MONTH, BOTH LINES



Δ 43% less downtime

|                          |                          |
|--------------------------|--------------------------|
| Warning before a failure | none → 9 days on average |
|--------------------------|--------------------------|

|             |                                   |
|-------------|-----------------------------------|
| Lost output | 504,000 bottles → 288,000 bottles |
|-------------|-----------------------------------|

|                            |        |
|----------------------------|--------|
| Emergency repair call-outs | 11 → 4 |
|----------------------------|--------|

|                                                  |             |      |
|--------------------------------------------------|-------------|------|
| GUARDRAIL Planned maintenance hours (limit +20%) | 20 h → 23 h | held |
|--------------------------------------------------|-------------|------|

### WHAT CHANGED

Vibration and temperature sensors on the motors, pumps and fillers. A model flags unusual patterns days ahead and opens a work order for a planned stop.

#### WHY THIS PROBLEM IS REAL

Augury, valued above \$1B, monitors machines for PepsiCo, Colgate-Palmolive and Hershey's. Siemens estimates unplanned downtime costs the world's 500 largest companies \$1.4 trillion a year, 11% of their revenue.

Fictional company. The figures are illustrative and anchored to the real benchmarks named on this tab.



SALES & MARKETING · FICTIONAL EXAMPLE

## Answering every inbound lead in minutes, before it goes cold

Industrial equipment distributor. 900 inbound leads a month from the website, phone and messaging apps.

PRIMARY METRIC · MEDIAN TIME TO FIRST REPLY



|                                               |                |
|-----------------------------------------------|----------------|
| Leads never contacted                         | 21% → 0%       |
| Leads that become a meeting                   | 12% → 19%      |
| Meetings per month                            | 108 → 171      |
| GUARDRAIL Meetings sales accepts as qualified | 64% → 66% held |

### WHAT CHANGED

Every lead gets an instant reply that asks three qualifying questions and offers the right rep's calendar. Reps start the day with meetings, not a pile of forms.

#### WHY THIS PROBLEM IS REAL

Salesforce paid about \$1.2B for Qualified, whose AI SDR answers inbound leads instantly. A Harvard Business Review study of 2,241 companies found that firms replying within an hour were about 7× more likely to qualify a lead. The average reply took 42 hours.

Fictional company. The figures are illustrative and anchored to the real benchmarks named on this tab.



FINANCE & ACCOUNTING · FICTIONAL EXAMPLE

## Closing the month in 4 days instead of 9

Holding company with six subsidiaries and one consolidated monthly close.

PRIMARY METRIC · BUSINESS DAYS TO CLOSE THE MONTH



|                                             |                |
|---------------------------------------------|----------------|
| Hours on reconciliations per month          | 160 h → 55 h   |
| Reconciliations done by hand in Excel       | 140 → 30       |
| Management report arrives                   | day 12 → day 6 |
| GUARDRAIL Adjustments found after the close | 7 → 3 held     |

### WHAT CHANGED

Bank, intercompany and account reconciliations are matched automatically, and only the exceptions reach an accountant. A shared checklist shows what is holding up the close.

#### WHY THIS PROBLEM IS REAL

FloQast, valued at \$1.6B, serves more than 2,600 accounting teams. APQC's benchmark of 2,300 organizations puts the median close at 6.4 days, and a 2025 survey found 94% of finance teams still reconcile in Excel.

Fictional company. The figures are illustrative and anchored to the real benchmarks named on this tab.



SUPPLY CHAIN · FICTIONAL EXAMPLE

## Fewer empty shelves with less stock

Food distributor. 1,800 products delivered to 40 stores and customers.

PRIMARY METRIC · ORDER LINES OUT OF STOCK



Δ stockouts halved

|                           |           |
|---------------------------|-----------|
| Forecast error (weighted) | 38% → 24% |
|---------------------------|-----------|

|                           |         |
|---------------------------|---------|
| Days of inventory on hand | 46 → 37 |
|---------------------------|---------|

|                                  |             |
|----------------------------------|-------------|
| Planner hours on orders per week | 60 h → 22 h |
|----------------------------------|-------------|

|                                     |             |      |
|-------------------------------------|-------------|------|
| GUARDRAIL Fresh product written off | 2.1% → 1.6% | held |
|-------------------------------------|-------------|------|

### WHAT CHANGED

A forecast for every product and location that accounts for promotions, seasons and holidays, turned into suggested orders. Planners review the exceptions instead of building every order.

#### WHY THIS PROBLEM IS REAL

RELEX Solutions, valued at \$5.8B, runs forecasting and replenishment for 350+ retailers and suppliers. IHL Group puts the global cost of out-of-stocks and overstocks at \$1.77 trillion a year, 6.5% of retail sales.

Fictional company. The figures are illustrative and anchored to the real benchmarks named on this tab.



PROCUREMENT · FICTIONAL EXAMPLE

## From request to purchase order in 4 days, not 16

Shared procurement desk for the whole group. 650 purchase requests a month.

PRIMARY METRIC · MEDIAN DAYS FROM REQUEST TO PURCHASE ORDER

|        |             |         |
|--------|-------------|---------|
| Before | <div></div> | 16 days |
| After  | <div></div> | 4 days  |

Δ 4× faster

|                                            |          |
|--------------------------------------------|----------|
| Requests sent back for missing information | 34% → 8% |
|--------------------------------------------|----------|

|                                   |       |
|-----------------------------------|-------|
| Manual approval steps per request | 5 → 2 |
|-----------------------------------|-------|

|                                      |              |
|--------------------------------------|--------------|
| Buyer hours spent chasing, per month | 210 h → 70 h |
|--------------------------------------|--------------|

|                                                |           |      |
|------------------------------------------------|-----------|------|
| GUARDRAIL Spend that follows purchasing policy | 71% → 90% | held |
|------------------------------------------------|-----------|------|

### WHAT CHANGED

One request form that asks the right questions for each category, routes by amount and risk, and approves small routine purchases automatically within policy.

#### WHY THIS PROBLEM IS REAL

Zip, valued at \$2.2B, runs purchasing for large enterprises and reports more than \$6B in customer savings. Dollar Tree cut its procurement cycle time by 70%.

Fictional company. The figures are illustrative and anchored to the real benchmarks named on this tab.

### METHOD

## Five stages. Each one ends at a gate.

## 01

**Find**

FIND THE PROBLEM

Walk the process with the people who run it. Pull timestamps from the systems.

Rank bottlenecks by volume × minutes per unit × cost of an hour.

WE DO

Check each problem twice: market proof (a funded startup solves it) and local proof (it costs this organization measurable time or money).

YOU GET

Bottleneck Map: every step, its time, and where work waits.

GATE

A named owner who wants it fixed, and roughly 40+ staff-hours a month at stake.

## 02

**Baseline**

MEASURE BEFORE

Agree exactly what we measure, then measure the process as it runs today.

WE DO

Metric Contract: one primary metric, 2–4 guardrails, start and stop events, data source. At least 30 units or 2 normal weeks.

YOU GET

Baseline Certificate with median, P90 and evidence level.

GATE

The process owner signs. From here nobody moves the goalposts, including us.

## 03

**Build**

TOOL AND SOLUTION

The smallest intervention that moves the measured number. Sometimes that is AI, sometimes a form or a changed hand-off.

WE DO

Human in the loop wherever judgment or risk is involved. Shadow mode first: the tool runs while people still do the work.

YOU GET

Working tool, an exception path for cases it can't handle, and a run book.

GATE

Handles a real sample of cases at acceptable quality in shadow mode.

# 04

## Prove

MEASURE AFTER

Same metric, same definition, same data source as the baseline.

WE DO

Count all human time, including review, corrections and exceptions. Run a control group in the same period where possible.

YOU GET

Impact Report: before, after,  $\Delta$ , evidence level, guardrails, and what did not work.

GATE

Target met and no guardrail worse. If not, back to Build, and the report says so.

# 05

## Scale

EFFICIENCY, THEN REPEAT

Roll out to full volume, keep the metric live so drift shows, and add the result to the Value Ledger.

WE DO

Rollout, drift monitoring, owner training, monthly ledger update.

YOU GET

A ledger entry leadership can read in one line, and the next bottleneck.

GATE

Monitoring live and the next bottleneck selected.



Fixing one constraint moves the constraint somewhere else. Stage 05 hands the next bottleneck back to Stage 01.

### — MEASURES

## Efficiency is four numbers, not one.

Every project picks one primary delta and protects the others with guardrails.

### Time

How fast does one unit of work get through?

min / unit

lead time, h

wait, days

### Quality

How often is it right the first time?

error rate

rework %

hit rate

## Cost

What does one unit cost, all-in?

cost / unit

staff-h / month

tool cost

## Capacity

How much more can the same team handle?

units / person-day

backlog

peak volume

## Guardrails: what must not get worse

A resume screen that is 20× faster because it rejects everyone fails its quality guardrail. Speed from cutting corners is not a result.

ranking fairness

candidate experience

employee satisfaction


compliance

downstream workload

### EVIDENCE


## Every number we report carries its evidence level.

**L1 Estimate** Interviews: "it takes about five minutes."




Used to prioritize. Never reported as a result.

**L2 Observed** A timed sample by us, 30 or more units.



Acceptable baseline when the process leaves no logs.

**L3 Logged** System timestamps for the full population.



Preferred for both baseline and result.

**L4 Controlled** Parallel run or control group in the same period.



Required whenever our fee depends on the result.

People work faster when they know they are being timed. That is why system logs beat stopwatches whenever both exist.

## Eight rules we don't bend.

**R1 No baseline, no build.**

No code is written before the baseline is signed.

**R3 Same ruler before and after.**

Same definition, same data source, same start and stop events.

**R5 One primary metric, a few guardrails.**

A project that improves everything at once proves nothing.

**R7 Report what did not work.**

Failed experiments go into the report. That is what makes the successes believable.

**R2 Measure the work, not the tool.**

"Minutes per resume from arrival to a complete ATS record", not "the model answers in two seconds".

**R4 Count all the human time.**

Review, correction and exception handling belong to the after number.

**R6 The owner signs, not us.**

The process owner signs both the baseline and the result.

**R8 Hours released are capacity, not cash.**

They become savings when the capacity is redeployed or a hire is avoided. We report both and say which happened.

## Resume intake into the ATS

Built · baseline in progress

Resumes arrive by email and job boards, and an operator types each one into the ATS. Recruiters only see a candidate after that step, so every day of backlog is a day good candidates wait. The tool creates the record, the operator reviews it, and a ranking decides which candidates recruiters open first. The ranking sets the order of review. It never rejects anyone.

| ROLE      | METRIC               | DEFINITION                                                                                             | SOURCE                                 |
|-----------|----------------------|--------------------------------------------------------------------------------------------------------|----------------------------------------|
| PRIMARY   | Handling time        | Operator minutes per resume, from opening it to a complete ATS record, including review and correction | Timed sample (L2), ATS timestamps (L3) |
| SECONDARY | Time to first review | Hours from resume arrival until a recruiter first opens the record                                     | Email arrival + ATS view log           |
| GUARDRAIL | Field accuracy       | Share of fields correct against the source resume, 50 random records                                   | Record audit                           |
| GUARDRAIL | Shortlist hit rate   | Share of candidates later invited to interview who were in the tool's top 20%                          | ATS outcomes, 60 days                  |
| GUARDRAIL | Ranking fairness     | Each group's top-20% rate is at least four-fifths of the highest group's rate                          | Audit, same 60 days                    |

## INPUTS · EXAMPLE VALUES

Resumes per month

|      |       |
|------|-------|
| 1200 | units |
|------|-------|

Handling time before

|   |     |
|---|-----|
| 6 | min |
|---|-----|

Handling time after, incl. review

|     |     |
|-----|-----|
| 0.3 | min |
|-----|-----|

Loaded cost of one staff hour

|    |     |
|----|-----|
| 15 | USD |
|----|-----|

One-time build cost

|      |     |
|------|-----|
| 3000 | USD |
|------|-----|

Running cost per month

|     |     |
|-----|-----|
| 100 | USD |
|-----|-----|

Currency label

|     |
|-----|
| USD |
|-----|

Evidence level of these inputs: L1. Replace them with the signed baseline and the measured after.

# 20×

6:00 → 0:18 per resume

Staff-hours released / month

**114 h**

Full-time equivalent (160 h)

**0.71**

Net value / month

**1,610 USD**

Payback

**1.9 months**

STAFF-HOURS PER MONTH



Net value over 12 months: 19,320 USD. Hours released are capacity. They become savings when that capacity is redeployed or a hire is avoided.

## — LEDGER

# The Value Ledger only counts what has been measured.

One page for leadership: every project, its primary metric, its proven  $\Delta$  and the evidence behind it. This is the ledger on day one. Case 001 enters the total only after its after-measurement is signed.

| #   | PROCESS                    | PRIMARY METRIC                      | PROVEN $\Delta$ | EVIDENCE | STATUS    |
|-----|----------------------------|-------------------------------------|-----------------|----------|-----------|
| 001 | Resume intake into the ATS | Minutes per resume                  | not yet         | L1       | Baseline  |
| 002 | Interview scheduling       | Hours from invite to confirmed slot | not yet         | none     | Candidate |
| 003 | HR question desk           | First-response time                 | not yet         | none     | Candidate |
| 004 | Monthly HR report          | Staff-hours to produce              | not yet         | none     | Candidate |

Proven so far: 0 staff-hours / month

— WORK WITH US

# Three steps, each one small enough to say no to.

STEP 1 · 2 WEEKS

## Bottleneck Audit

We walk your processes, pull the logs, and rank where time and money leak.

YOU GET

Bottleneck Map, top three opportunities ranked by expected  $\Delta$ , draft Metric Contracts.

PRICING

No fee. The audit is our investment in finding a difference worth making.

STEP 2 · 4–6 WEEKS

## Proof Sprint

One bottleneck, taken from signed baseline to measured result.

YOU GET

Baseline Certificate, working tool, shadow run, Impact Report.

PRICING

No measurable difference, no fee. Target met: the standard fee agreed before the baseline.

STEP 3 · ONGOING

## Scale & Ledger

Rollout, drift monitoring, and the next bottleneck.

YOU GET

Monthly Value Ledger and a steady queue of proven improvements.

PRICING

Standard plan or Ambitious plan, below.

## Pay only for the difference

STANDARD PLAN

For teams that want one clear, proven improvement.

- Measured difference below the agreed target: **you pay nothing.**
- Target met: **one standard fee**, written into the contract before we measure anything.

## AMBITIOUS PLAN

For teams that want a major productivity program. The fee rises in fixed, pre-agreed steps as the measured improvement grows.

|                                                                                                                                                                                                                                                                    |  |      |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--|------|
| Below the target                                                                                                                                                                                                                                                   |  | 0    |
| A difference counts only when it is measured at evidence level L3 or L4, with the same method as the signed baseline, and no guardrail got worse. Intercom, Sierra and HighRadius already charge per measured outcome. We apply the same logic to whole processes. |  |      |
| Target reached                                                                                                                                                                                                                                                     |  | 1.0× |
| 1.5× the target                                                                                                                                                                                                                                                    |  | 1.4× |
| 2× the target                                                                                                                                                                                                                                                      |  | 1.8× |
| 3× the target or more                                                                                                                                                                                                                                              |  | 2.5× |
| Inside a holding company the same steps run as an internal Efficiency Lab; each subsidiary is a client, and the Value Ledger goes to leadership every month.                                                                                                       |  | cap  |

Fee shown as a multiple of the standard fee. Example: the target is to close the month 3 days sooner. Closing it 4.5 days sooner is 1.5× the target, so the fee is 1.4× the standard fee.

## OUR PROMISE

**We find the bottleneck. We measure it. We fix it. Then we measure it again, with your data and your signature.**

We agree the target after the baseline. If the number doesn't move, you pay nothing.

CONTACT

# Start with a free Bottleneck Audit.

Tell us which process feels slow. We measure it, and you decide with the numbers in front of you.



WEBSITE

**deltaup.ir**



PHONE

**0912 836 1573**



SOCIAL MEDIA

**0930 695 2808**

Find us and message us on this number.

DeltaUp · Delta Method v1    deltaup.ir · 0912 836 1573

Problem source: HR problem atlas, 483 startups across 14 HR domains

All Case 001 figures are examples until the baseline is signed